

SegNet implementation in PyTorch

Move the frame to colab [Done]

Mount data [done]

Fix the training function [Done]

Report training accuracy [Done]

Learning Rate decay and cooldown [Done]

Save checkpoints [Done]

Visualize an output [Done]

Plot training curve [Done]

Calculate class weights median frequency balancing [Done]

Online Data Augmentation (intensity, noise, stretch, center shift) [Done]

Quantitive evaluation (accuracy, etc) [Done]

Video visualization [Done]

Demo on CamVid Seq01, Kitti(preprocessing?) [Done]

In [1]:

```
# import necessary libraries
%matplotlib inline
from matplotlib import pyplot as plt
import torch
from torchvision import models
import torchsummary
import torch.nn as nn
import torch.nn.functional as F
import cv2
import numpy as np
```

```
import numpy as np
from torch.utils.data import DataLoader
import time
import os

from tqdm import tqdm_notebook as tqdm
```

In [2]:

```
# Mount Google drive
from google.colab import drive
drive.mount('/content/gdrive')

# !ls gdrive/MyDrive
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

In []:

```
# get the path to the camvid data
camvid_path = r'/content/gdrive/MyDrive/CamVid'
!ls '/content/gdrive/MyDrive/CamVid'

# get the path to the scripts
!ls gdrive/MyDrive/Segmentation/code

# update imgaug
!pip uninstall imgaug
!pip install imgaug==0.4.0

# access necessary code
from gdrive.MyDrive.Segmentation.code.model import SegNet
import gdrive.MyDrive.Segmentation.code.data_aug as data
```

Global vars

In [4]:

```
# for reproduction
```

```

torch.manual_seed(123)

# Pre config
IN_CHANNELS = 3
NUM_CLASS = 2

# training parameters
NUM_EPOCHS = 200
LEARNING_RATE = 0.1 # will decrease with epoch growing larger
BATCH_SIZE = 16
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
LABEL_WEIGHTS = torch.FloatTensor([0.5, 2])
WEIGHT_DECAY = 0.0005
NUM_WORKERS = 0
PIN_MEMORY = False

```

Reporting/Evaluation Functions

In [5]:

```

def avg_pixelwise_accuracy(model, dataset):

    if not dataset:
        return -1

    # change to eval mode
    model.eval()

    # get the dataloader using large batch since only forward pass required
    loader = DataLoader(dataset=dataset,
                        batch_size=8,
                        )

    correct, total = 0, 0
    for i, batch in enumerate(loader):

        # mount to GPU if available [Need to fix y]
        imgs, labels = batch['X'].to(DEVICE), batch['Y'].to(DEVICE)

        # argmax X along dim 1
        out = model(imgs)
        predicted = torch.argmax(out, dim=1)
        total += labels.nelement()
        correct += predicted.eq(labels).sum().item()

    return correct / total

```

Save/Load checkpoints

In [6]:

```
def save_checkpoints(model, filename, epoch, optimizer, loss, train_acc, valid_acc):

    dirname = "/content/gdrive/MyDrive/Segmentation/check_points/large_aug"
    os.makedirs(dirname, exist_ok=True)

    save_path = os.path.join(dirname, filename)
    print(save_path)
    torch.save({
        'epoch': epoch,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'loss': loss,
        'train_acc': train_acc,
        'valid_acc': valid_acc,
        'transfer_learning': TRANSFER_LEARNING
    }, save_path)

def load_checkpoints(filename, transfer_learning=True):

    dirname = "/content/gdrive/MyDrive/Segmentation/check_points/large_aug"
    file_path = os.path.join(dirname, filename)
    checkpoint = torch.load(file_path)
    if 'transfer_learning' in checkpoint:
        transfer_learning = checkpoint['transfer_learning']
    model = SegNet(IN_CHANNELS, NUM_CLASS, transfer_learning=transfer_learning)
    model.load_state_dict(checkpoint['model_state_dict'])

    print("Model Loaded: Epoch#{0}\tLoss:{1:.4f}\tTrainAcc:{2:.4f}\tValidAcc:{3:.4f}".format(
        checkpoint['epoch'],
        checkpoint['loss'],
        checkpoint['train_acc'],
        checkpoint['valid_acc'],
    ))
    return model.to(DEVICE)

# model2 = load_checkpoints("Epoch1_loss59.1495_trainacc31.217_valacc28.866.pth")
# model2
```

Visualization

In [7]:

```
def show_img(img_tensor):

    if len(img_tensor.shape) == 4:
        img = img_tensor.squeeze()
    else:
        img = img_tensor
    img = torch.movedim(img, 0, -1).type(torch.uint8).numpy()[:, :, ::-1]
    plt.imshow(img)
    plt.show()

def show_mask(mask_tensor):
    # shape height x width, single channel
    plt.imshow(mask_tensor)
    plt.show()

def show_pred_mask(model, img_tensor):
    if len(img_tensor.shape) == 4:
        source_img = img_tensor
    else:
        source_img = img_tensor.unsqueeze(0).to(DEVICE)

    model.eval()
    out = model(source_img)
    predicted = torch.argmax(out, dim=1)
    predicted = predicted.cpu().data.numpy()[0, :]
    plt.imshow(predicted)
    plt.show()

def show_all(model, img_tensor, mask_tensor, cmap='jet'):

    # regulate the image tensor
    if len(img_tensor.shape) == 4:
        img = img_tensor.squeeze()
        source_img = img_tensor
    else:
        img = img_tensor
        source_img = img_tensor.unsqueeze(0).to(DEVICE)

    # Raw image
    img = torch.movedim(img, 0, -1).type(torch.uint8).numpy()[:, :, ::-1]

    # Ground Truth mask
    mask = mask_tensor

    # predicted mask
    model.eval()
    out = model(source_img)
    predicted = torch.argmax(out, dim=1)
    predicted = predicted.cpu().data.numpy()[0, :]
```

```

fig, axs = plt.subplots(1, 5, figsize=(20, 20), constrained_layout=True)

axs[0].imshow(img)
axs[0].set_title("Image")

axs[1].imshow(mask)
axs[1].set_title("Ground Truth")

axs[2].imshow(predicted)
axs[2].set_title("SegNet Outcome")

img[:, :, 1] += (predicted * 255 * 0.15).astype(np.uint8)
img[:, :, 0] += (predicted * 255 * 0.3).astype(np.uint8)
axs[3].imshow(img, vmin=0, vmax=255)
# axs[3].imshow(predicted, cmap=cmap, alpha=0.3)
axs[3].set_title("SegNet Overlay")

axs[4].imshow((predicted - mask.numpy()) ** 2)
axs[4].set_title("Error")

plt.show()

```

Dataset Preparation

In [8]:

```

# load the dataset
X_train, y_train = data.load_data(camvid_path+"/train", im_height=360, im_width=480, aug_copies=5)
print("Loaded {} training samples".format(len(X_train)))

X_valid, y_valid = data.load_data(camvid_path+"/val", im_height=360, im_width=480)
print("Loaded {} validation samples".format(len(X_valid)))

X_test, y_test = data.load_data(camvid_path+"/test", im_height=360, im_width=480)
print("Loaded {} testing samples".format(len(X_test)))

```

Bkg : Lane = 43570264.0 : 19847338.0
 Loaded 2202 training samples

Bkg : Lane = 12400839.0 : 5051963.0
 Loaded 101 validation samples

Bkg : Lane = 29962033.0 : 10300369.0
 Loaded 233 testing samples

Train

In [9]:

```
# Observe an image from Kitti
kitti_observe = cv2.imread("/content/gdrive/MyDrive/KiTTi_dataset/testing_crop/um_000052.png")
kitti_observe_tensor = torch.from_numpy(np.moveaxis(kitti_observe, -1, 0))
```

In [10]:

```
# train function
from torch.utils.data import Dataset
class CamVidDataset(Dataset):
    def __init__(self, imgs, masks):
        self.X = imgs
        self.y = masks
    def __len__(self):
        return len(self.X[:, 0, 0, 0])

    def __getitem__(self, index):
        data = {'X': torch.FloatTensor(self.X[index, :, :, :]), 'y': torch.LongTensor(self.y[index, :, :, :])}
        return data

def train(model, optimizer, loss_fn, X_train, y_train, X_valid=None, y_valid=None):

    # get the dataloader
    train_dataset = CamVidDataset(X_train, y_train)
    train_loader = DataLoader(dataset=train_dataset,
                              batch_size=BATCH_SIZE,
                              num_workers=NUM_WORKERS,
                              pin_memory=False,
                              shuffle=True,
                              drop_last=True,
                              worker_init_fn=None
                             )

    valid_dataset = CamVidDataset(X_valid, y_valid) if X_valid is not None else None

    # for plotting and logging
    epoch_lst, iters, losses, train_acc_lst, val_acc_lst, iter_counter = [], [], [], [], [], 0

    # for checkpoints
    checkpoint_path_template = "Epoch{}_loss{:.4f}_trainacc{:.3f}_valacc{:.3f}.pth"
    acc_recorder, save_gap = 96, 0.2

    # start training
    for epoch in range(NUM_EPOCHS):

        # visualize an segmentation every 10 epoch
        if epoch % 4 == 0:
            show_all(model, X_valid[0, :, :], y_valid[0, :, :])
            show_all(model, kitti_observe_tensor.float(), kitti_observe_tensor.float()[0, :, :])
```

```

# Learning rate decay [optional]
if (epoch + 1) % 15 == 0:
    optimizer.param_groups[0]['lr'] /= 3
    print("<===== Learning Rate {} -> {}===== >".format(optimizer.param_groups[0]['lr'], epoch))

epoch_loss = 0
t_start = time.time()
for i, batch in enumerate(train_loader):

    # mount to GPU if available [Need to fix y]
    imgs, labels = batch['x'].to(DEVICE), batch['y'][:, :, :, :].to(DEVICE)

    # change the mode to training mode and step training
    model.train()
    out = model(imgs)
    loss = loss_fn(out, labels)    # note: soft-max should not be used here since it's included in nn.CrossEntropyLoss
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()

    # save the current training information, TODO
    losses.append(loss)
    epoch_loss += loss
    iters.append(iter_counter)
    iter_counter += 1

delta = time.time() - t_start

# train_acc = avg_pixelwise_accuracy(model, train_dataset)
train_acc = avg_pixelwise_accuracy(model, train_dataset) * 100
valid_acc = avg_pixelwise_accuracy(model, valid_dataset) * 100
print("Epoch #{}\tLoss: {:.6f}\tTrain Acc: {:.3f}\tValid Acc: {:.3f}\tTime: {:.2f}s".format(epoch, epoch_loss, train_acc, valid_acc, delta))

train_acc_lst.append(train_acc)
val_acc_lst.append(valid_acc)
epoch_lst.append(epoch)

# Save checkpoints
if valid_acc > acc_recoder + save_gap:
    acc_recoder = valid_acc
    checkpoint_name = checkpoint_path_template.format(epoch+1, epoch_loss, train_acc, valid_acc)
    save_checkpoints(model, checkpoint_name, epoch+1, optimizer, epoch_loss, train_acc, valid_acc)

# Plot the curve
plt.title("Learning Curve")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Learning Curve")
plt.plot(epoch_lst, train_acc_lst, label="Train")
plt.plot(epoch_lst, val_acc_lst, label="Validation")

```

```

plt.xlabel("Epoch")
plt.ylabel("Pixelwise Accuracy")
plt.legend(loc='best')
plt.show()

```

Train from scratch

In [11]:

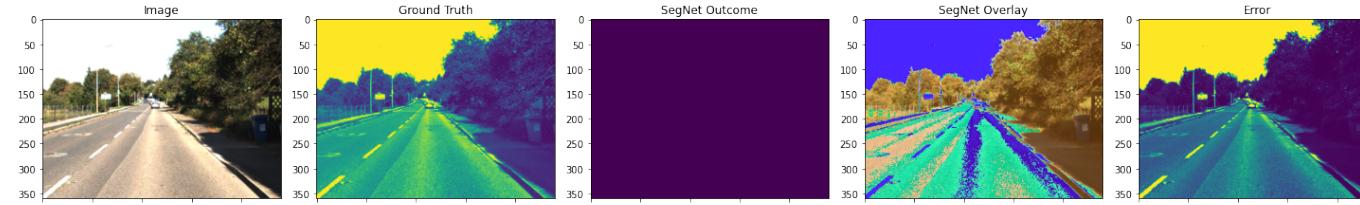
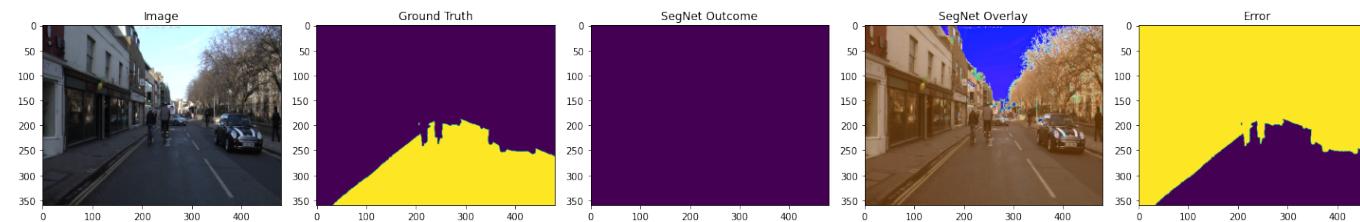
```

torch.cuda.empty_cache()

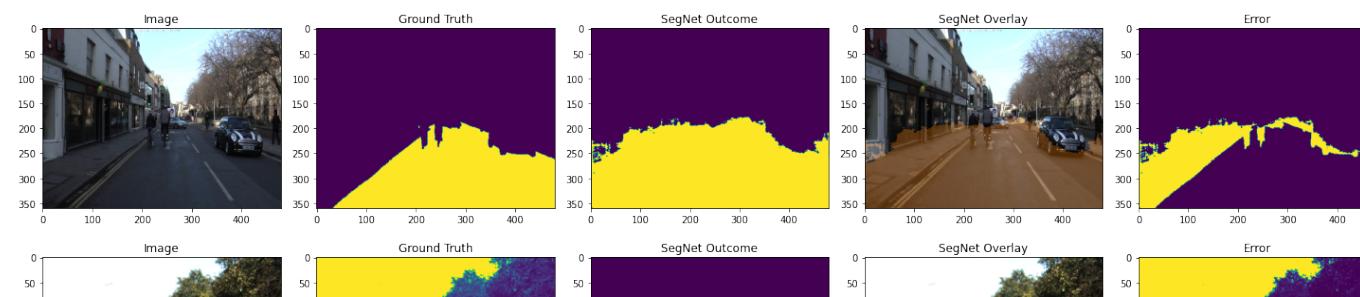
TRANSFER_LEARNING = False
model = SegNet(3,2, transfer_learning=TRANSFER_LEARNING).to(DEVICE)
optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)
loss = nn.CrossEntropyLoss(weight=LABEL_WEIGHTS.to(DEVICE))

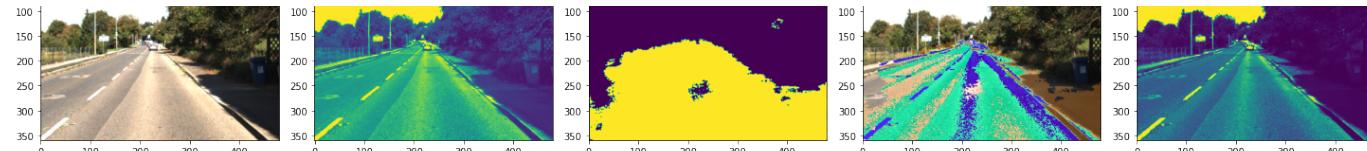
# Launch training
train(model, optimizer, loss, X_train, y_train, X_valid, y_valid)

```

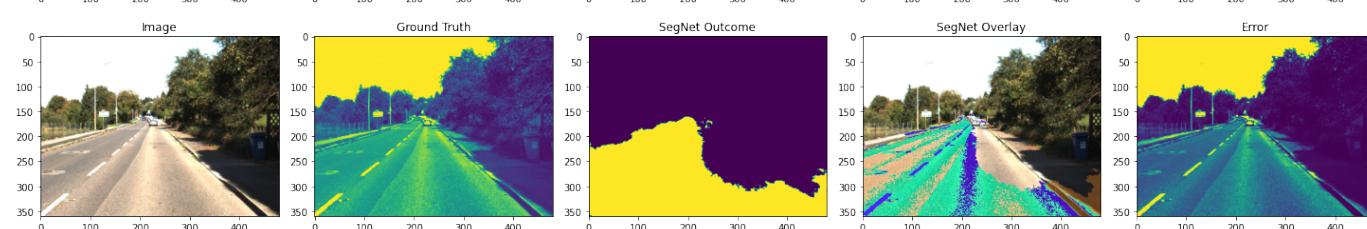
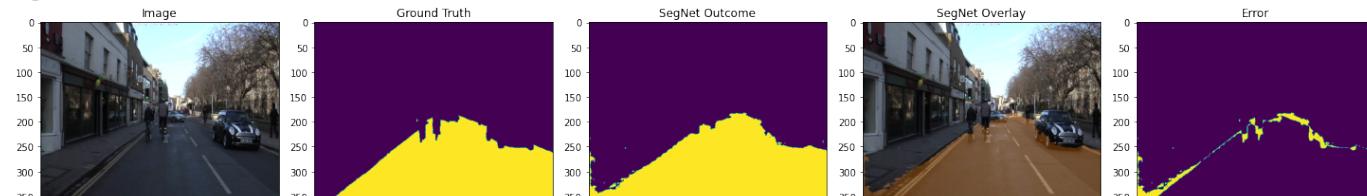


Epoch #1	Loss: 113.761894	Train Acc: 86.099%	Valid Acc: 88.964%	Time: 85.88s
Epoch #2	Loss: 41.867458	Train Acc: 90.610%	Valid Acc: 89.851%	Time: 85.89s
Epoch #3	Loss: 23.835253	Train Acc: 93.763%	Valid Acc: 93.365%	Time: 85.93s
Epoch #4	Loss: 19.350954	Train Acc: 90.274%	Valid Acc: 86.047%	Time: 85.89s

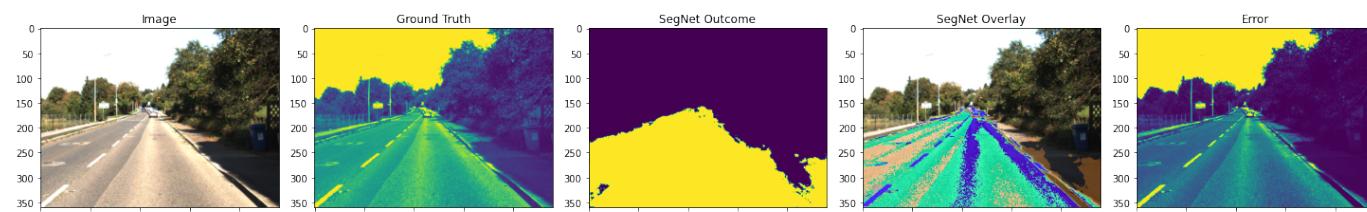
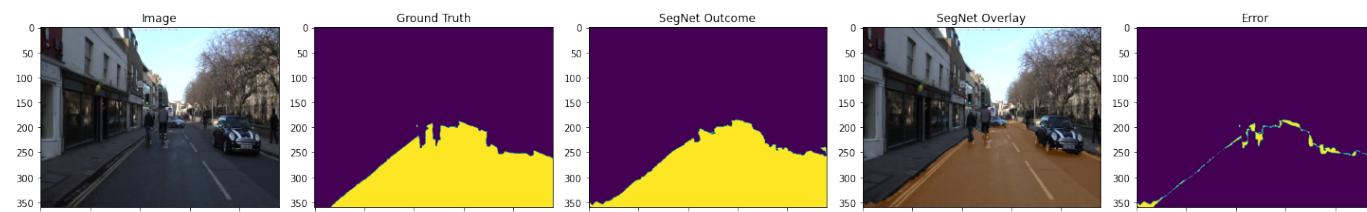




Epoch #5 Loss: 17.510023 Train Acc: 89.190% Valid Acc: 86.207% Time: 85.87s
 Epoch #6 Loss: 17.662569 Train Acc: 92.100% Valid Acc: 87.802% Time: 85.87s
 Epoch #7 Loss: 16.499863 Train Acc: 94.074% Valid Acc: 94.786% Time: 85.89s
 Epoch #8 Loss: 15.529844 Train Acc: 93.574% Valid Acc: 94.386% Time: 85.88s



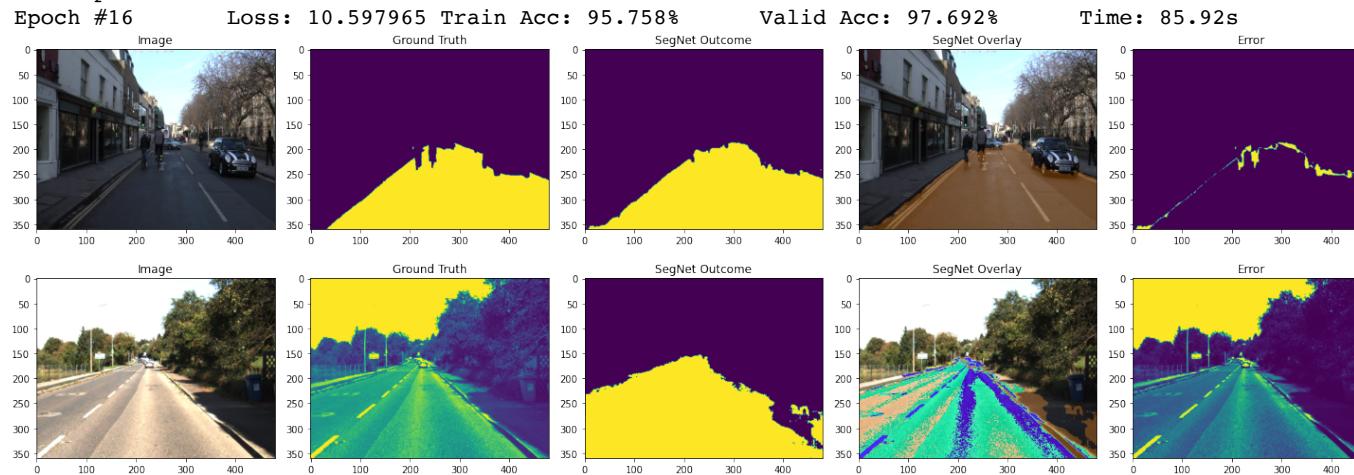
Epoch #9 Loss: 14.994182 Train Acc: 96.063% Valid Acc: 97.197% Time: 85.89s
 /content/gdrive/MyDrive/Segmentation/check_points/large_aug/Epoch9_loss14.9942_trainacc96.063_valacc97.197.pth
 Epoch #10 Loss: 14.117742 Train Acc: 94.361% Valid Acc: 95.133% Time: 85.86s
 Epoch #11 Loss: 13.702608 Train Acc: 95.265% Valid Acc: 96.897% Time: 85.90s
 Epoch #12 Loss: 13.000116 Train Acc: 95.437% Valid Acc: 97.075% Time: 85.89s



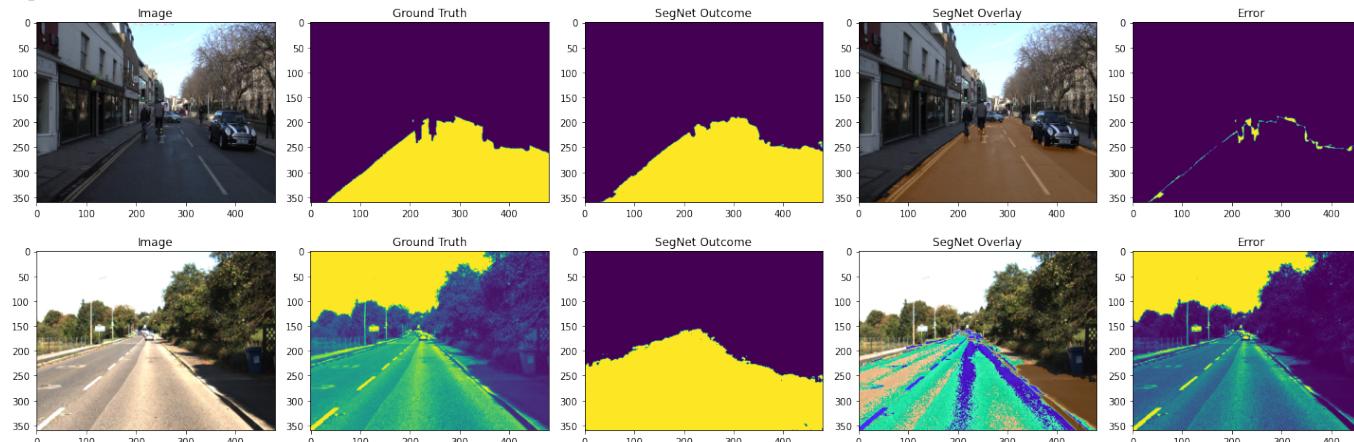
Epoch #13 Loss: 12.387345 Train Acc: 93.363% Valid Acc: 96.052% Time: 85.89s
 Epoch #14 Loss: 12.141721 Train Acc: 95.798% Valid Acc: 97.617% Time: 85.90s
 /content/gdrive/MyDrive/Segmentation/check_points/large_aug/Epoch14_loss12.1417_trainacc95.798_valacc97.617.pth

<===== Learning Rate 0.1 -> 0.03333333333333333 =====>

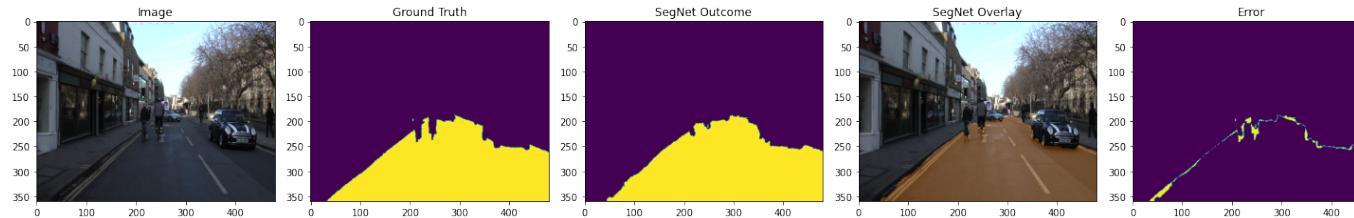
Epoch #15 Loss: 10.643561 Train Acc: 96.307% Valid Acc: 98.070% Time: 85.90s
/content/gdrive/MyDrive/Segmentation/check_points/large_aug/Epoch15_loss10.6436_trainacc96.307_valacc98.070.pth

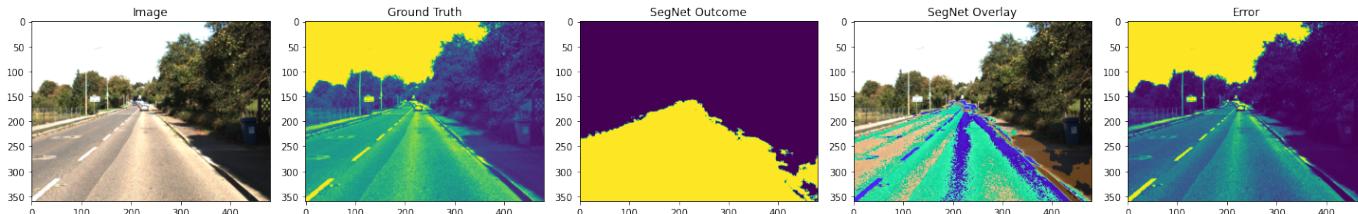


Epoch #17 Loss: 10.195443 Train Acc: 96.019% Valid Acc: 97.674% Time: 85.90s
Epoch #18 Loss: 9.994283 Train Acc: 96.044% Valid Acc: 97.943% Time: 85.90s
Epoch #19 Loss: 9.911497 Train Acc: 96.394% Valid Acc: 97.950% Time: 85.89s
Epoch #20 Loss: 9.835297 Train Acc: 95.538% Valid Acc: 97.824% Time: 85.93s

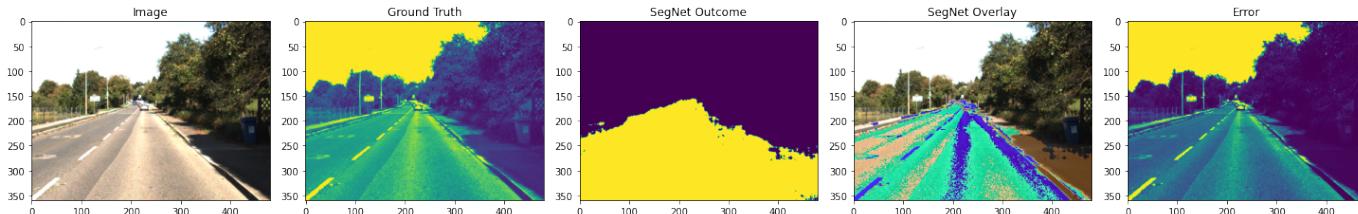
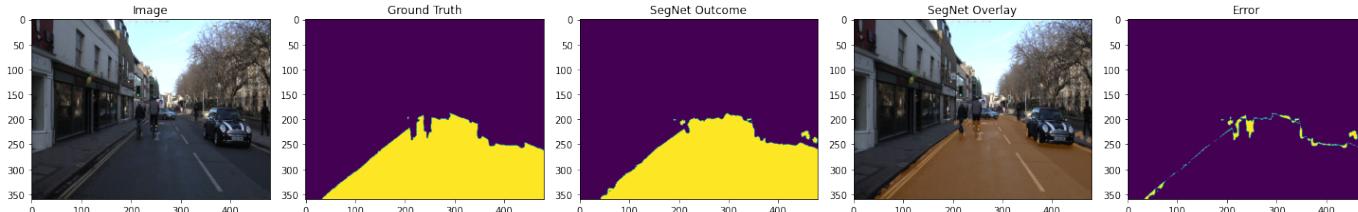


Epoch #21 Loss: 9.657845 Train Acc: 95.388% Valid Acc: 97.402% Time: 85.93s
Epoch #22 Loss: 9.526237 Train Acc: 95.580% Valid Acc: 97.826% Time: 85.95s
Epoch #23 Loss: 9.347882 Train Acc: 96.856% Valid Acc: 98.201% Time: 85.94s
Epoch #24 Loss: 9.458021 Train Acc: 96.388% Valid Acc: 98.263% Time: 85.89s





Epoch #25	Loss: 8.865716	Train Acc: 96.090%	Valid Acc: 98.095%	Time: 85.92s
Epoch #26	Loss: 8.797677	Train Acc: 96.723%	Valid Acc: 98.120%	Time: 85.95s
Epoch #27	Loss: 8.607855	Train Acc: 96.638%	Valid Acc: 97.870%	Time: 85.92s
Epoch #28	Loss: 8.652934	Train Acc: 96.506%	Valid Acc: 97.880%	Time: 85.93s



```
Epoch #29      Loss: 8.706271  Train Acc: 97.282%      Valid Acc: 98.383%      Time: 85.91s
/content/gdrive/MyDrive/Segmentation/check_points/large_aug/Epoch29_loss8.7063_trainacc97.282_valacc98
.383.pth
```

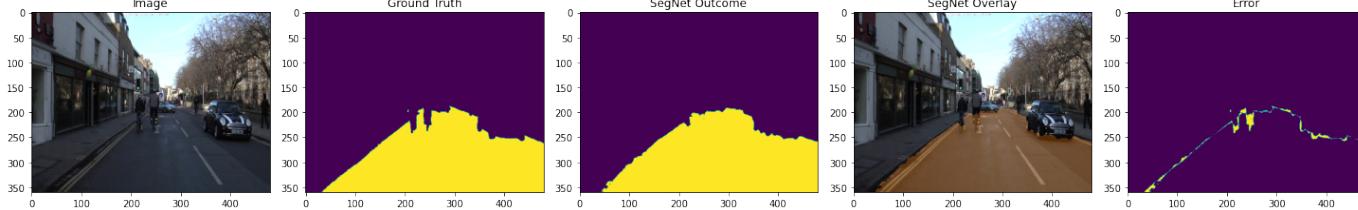
Epoch #30

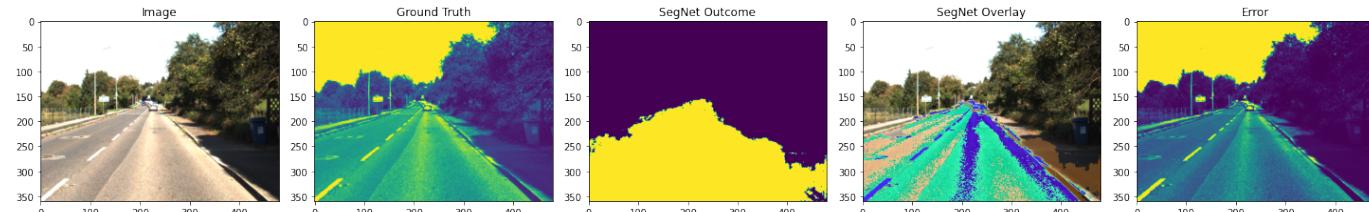
Epoch #31 Loss: 7.090017 Train Acc: 97.660% Valid Acc: 98.425% Time: 85.91

Epoch #32

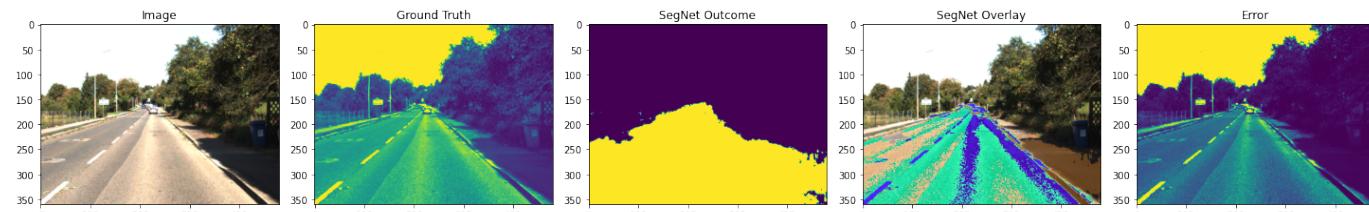
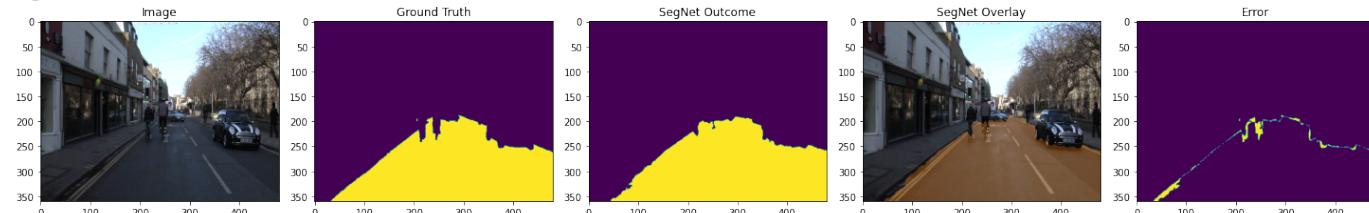
Image Ground Truth SegNet Outcome SegNet Overlay

50 50 50 50 50

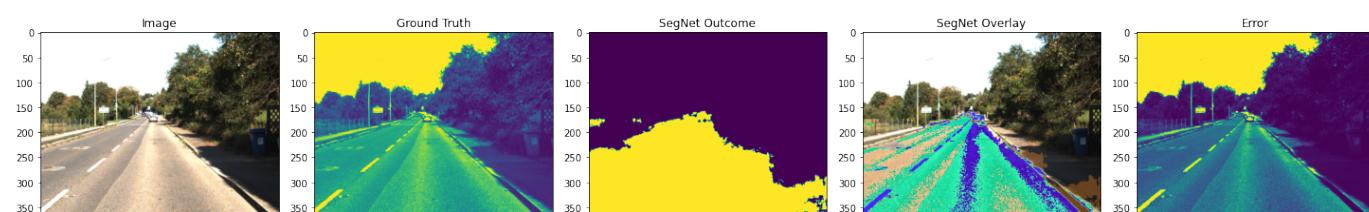
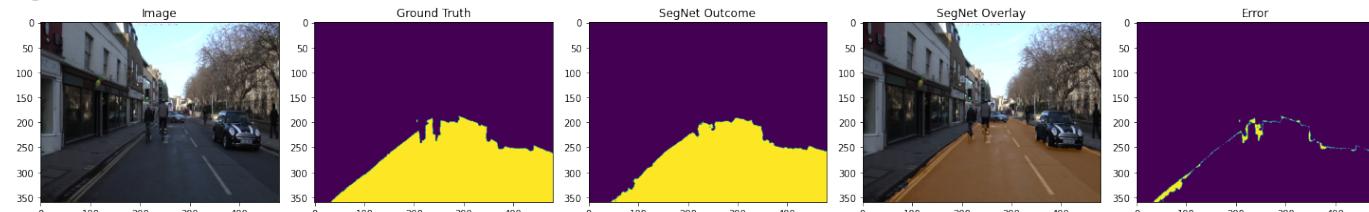




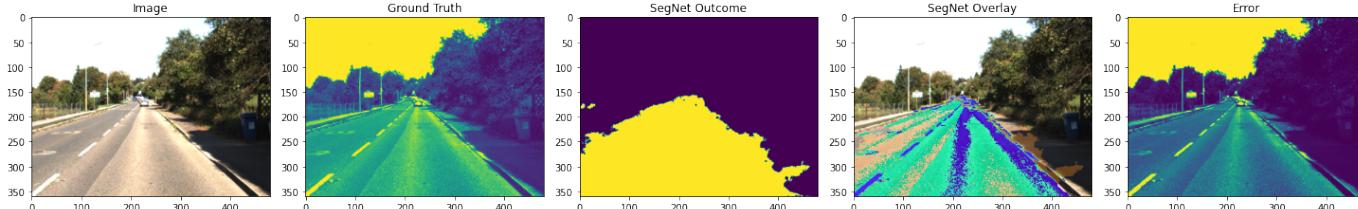
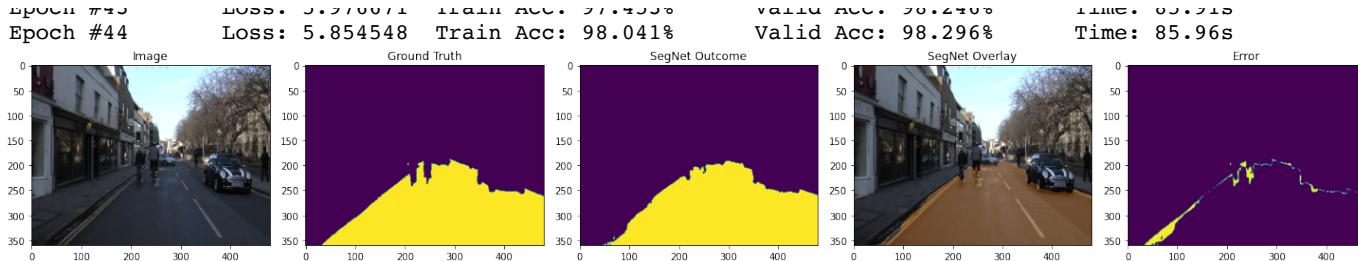
Epoch #33 Loss: 6.798064 Train Acc: 97.556% Valid Acc: 98.417% Time: 85.90s
Epoch #34 Loss: 6.549056 Train Acc: 97.632% Valid Acc: 98.395% Time: 85.91s
Epoch #35 Loss: 6.710833 Train Acc: 97.555% Valid Acc: 98.320% Time: 85.90s
Epoch #36 Loss: 6.522848 Train Acc: 97.153% Valid Acc: 98.134% Time: 85.90s



Epoch #37 Loss: 6.420828 Train Acc: 97.607% Valid Acc: 98.191% Time: 85.90s
Epoch #38 Loss: 6.245625 Train Acc: 97.786% Valid Acc: 98.446% Time: 85.92s
Epoch #39 Loss: 6.244062 Train Acc: 97.326% Valid Acc: 98.070% Time: 85.93s
Epoch #40 Loss: 6.117337 Train Acc: 97.685% Valid Acc: 98.306% Time: 85.91s



Epoch #41 Loss: 6.034916 Train Acc: 97.942% Valid Acc: 98.434% Time: 85.92s
Epoch #42 Loss: 5.891027 Train Acc: 97.912% Valid Acc: 98.394% Time: 85.94s
Epoch #43 Loss: 5.976671 Train Acc: 97.453% Valid Acc: 99.246% Time: 85.91s



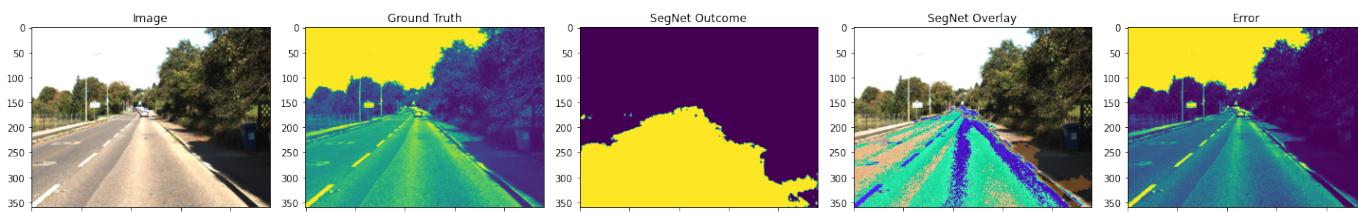
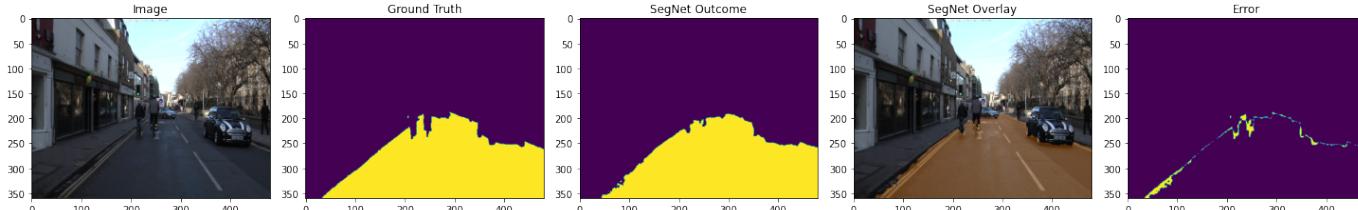
<===== Learning Rate 0.01111111111111112 -> 0.003703703703704 ===== =====>

Epoch #45 Loss: 5.181281 Train Acc: 98.107% Valid Acc: 98.373% Time: 85.90s

Epoch #46 Loss: 5.001117 Train Acc: 98.119% Valid Acc: 98.328% Time: 85.89s

Epoch #47 Loss: 4.907640 Train Acc: 98.031% Valid Acc: 98.322% Time: 85.91s

Epoch #48 Loss: 4.909502 Train Acc: 98.278% Valid Acc: 98.338% Time: 85.92s

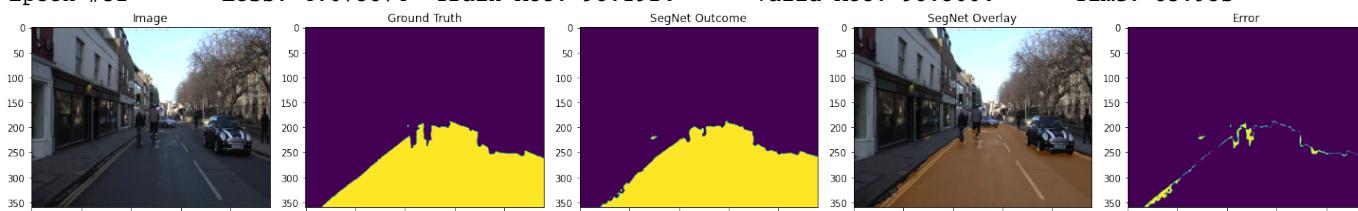


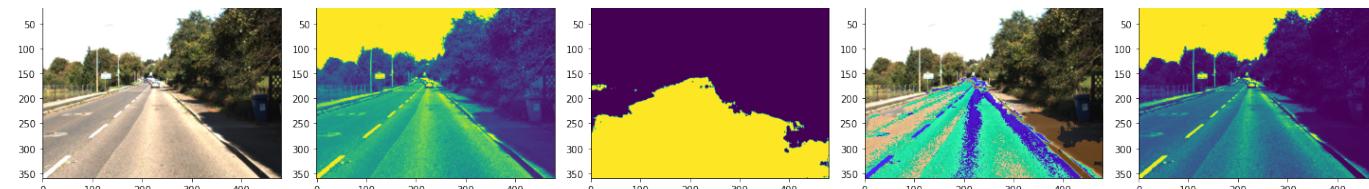
Epoch #49 Loss: 4.819682 Train Acc: 98.211% Valid Acc: 98.444% Time: 85.91s

Epoch #50 Loss: 4.781480 Train Acc: 98.278% Valid Acc: 98.485% Time: 85.94s

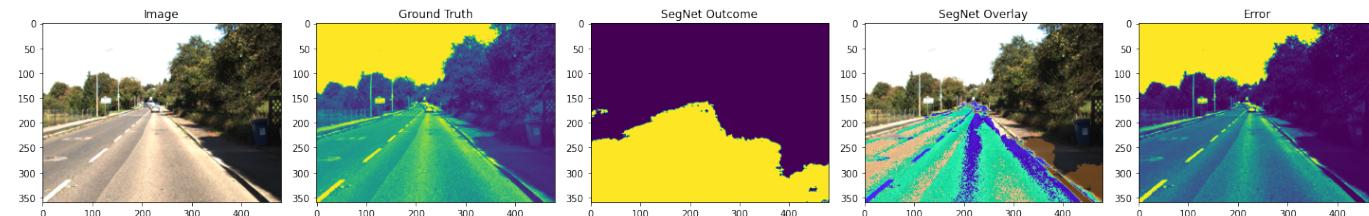
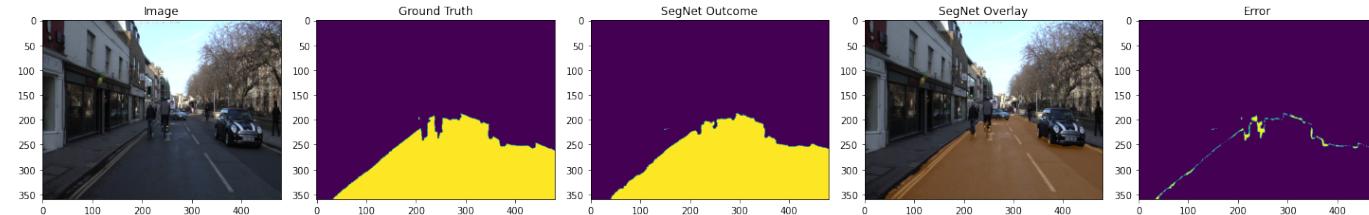
Epoch #51 Loss: 4.678722 Train Acc: 98.131% Valid Acc: 98.347% Time: 85.91s

Epoch #52 Loss: 4.673874 Train Acc: 98.192% Valid Acc: 98.366% Time: 85.93s

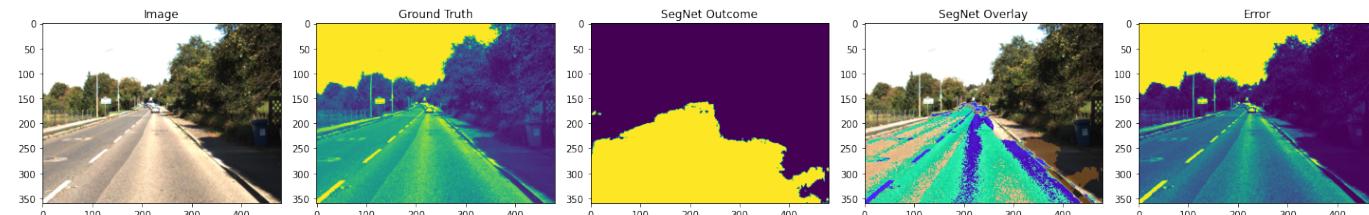
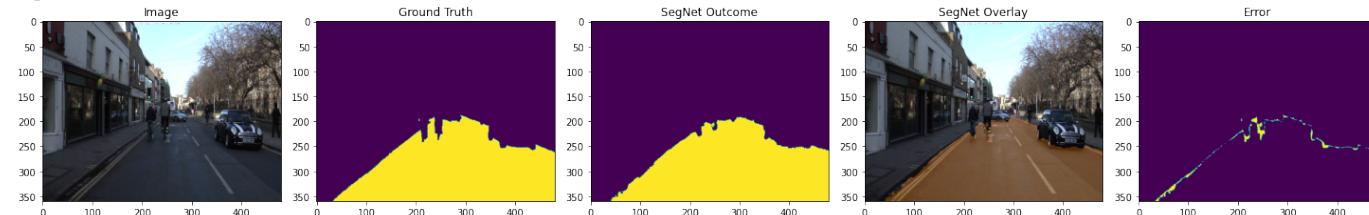




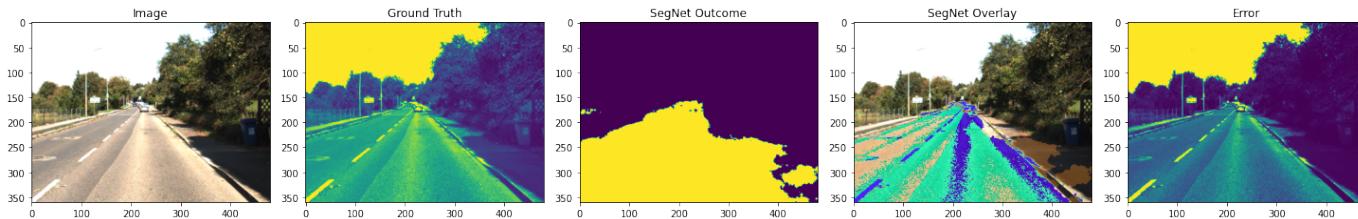
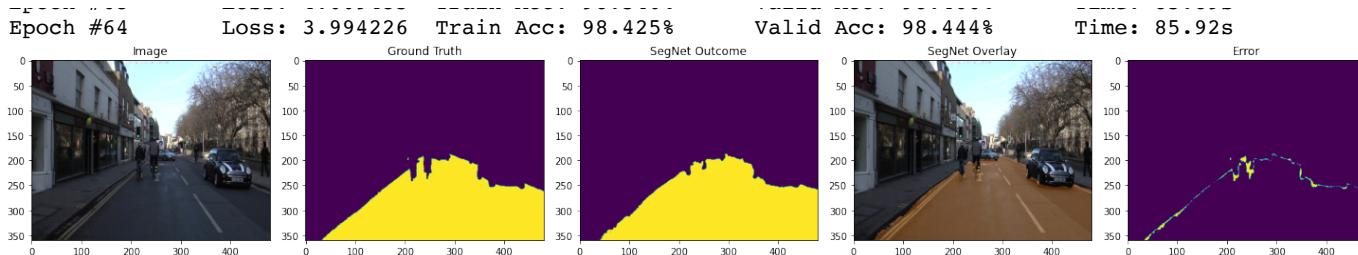
Epoch #53 Loss: 4.709645 Train Acc: 98.162% Valid Acc: 98.424% Time: 85.92s
 Epoch #54 Loss: 4.658604 Train Acc: 98.306% Valid Acc: 98.524% Time: 85.92s
 Epoch #55 Loss: 4.562550 Train Acc: 98.275% Valid Acc: 98.400% Time: 85.94s
 Epoch #56 Loss: 4.578988 Train Acc: 98.151% Valid Acc: 98.392% Time: 85.90s



Epoch #57 Loss: 4.528505 Train Acc: 98.461% Valid Acc: 98.482% Time: 85.92s
 Epoch #58 Loss: 4.430164 Train Acc: 98.352% Valid Acc: 98.381% Time: 85.92s
 Epoch #59 Loss: 4.453888 Train Acc: 98.374% Valid Acc: 98.417% Time: 85.91s
 <===== Learning Rate 0.003703703703704 -> 0.0012345679012345679 ===== =====>
 Epoch #60 Loss: 4.170705 Train Acc: 98.458% Valid Acc: 98.455% Time: 85.90s



Epoch #61 Loss: 4.075523 Train Acc: 98.354% Valid Acc: 98.390% Time: 85.89s
 Epoch #62 Loss: 4.035420 Train Acc: 98.425% Valid Acc: 98.395% Time: 85.90s
 Epoch #63 Loss: 4.009435 Train Acc: 98.540% Valid Acc: 98.466% Time: 85.89s



Epoch #65

Loss: 3.996300

Train Acc: 98.452%

Valid Acc: 98.417%

Time: 85.93s

Epoch #66

Loss: 3.921269

Train Acc: 98.517%

Valid Acc: 98.417%

Time: 85.88s

Epoch #67

Loss: 3.906095

Train Acc: 98.568%

Valid Acc: 98.451%

Time: 85.88s

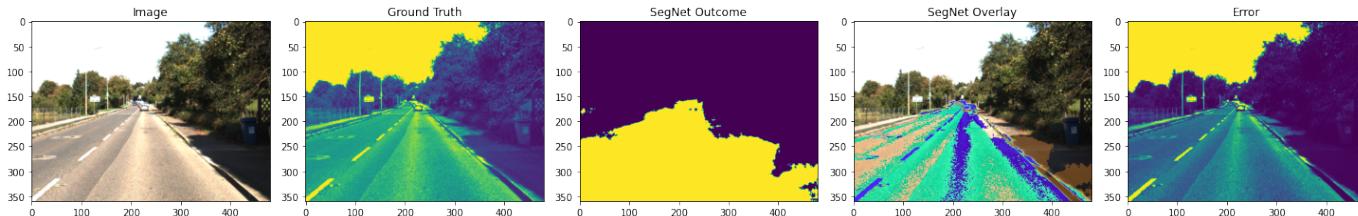
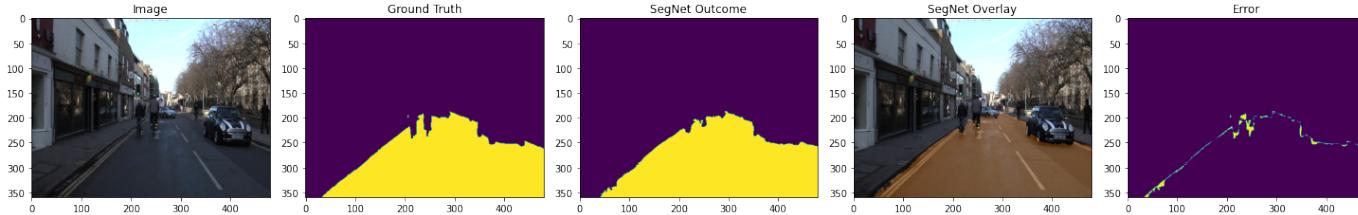
Epoch #68

Loss: 3.910293

Train Acc: 98.504%

Valid Acc: 98.450%

Time: 85.92s



Epoch #69

Loss: 3.928397

Train Acc: 98.541%

Valid Acc: 98.486%

Time: 85.86s

Epoch #70

Loss: 3.887527

Train Acc: 98.602%

Valid Acc: 98.455%

Time: 85.90s

Epoch #71

Loss: 3.856665

Train Acc: 98.422%

Valid Acc: 98.402%

Time: 85.94s

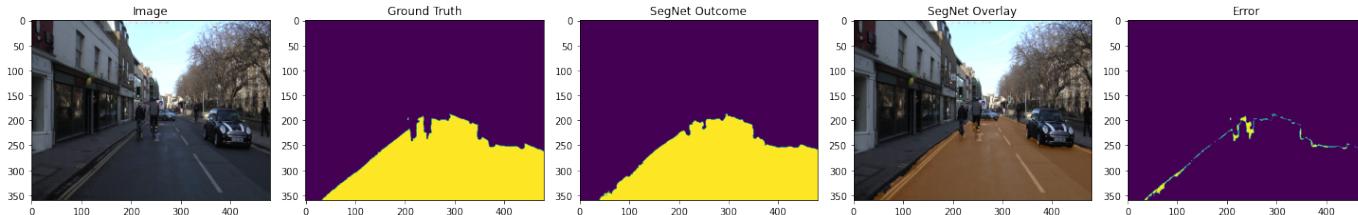
Epoch #72

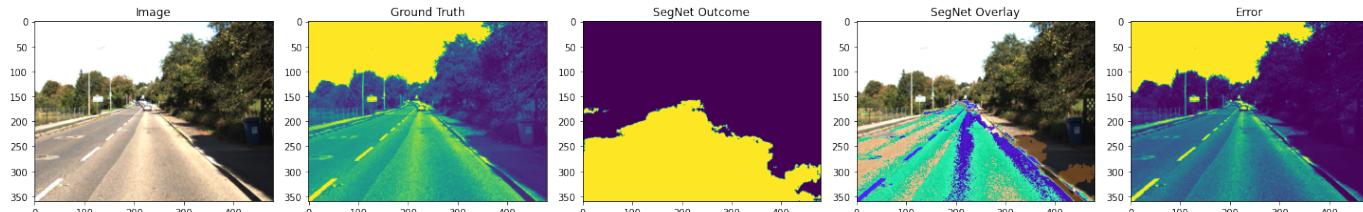
Loss: 3.861521

Train Acc: 98.540%

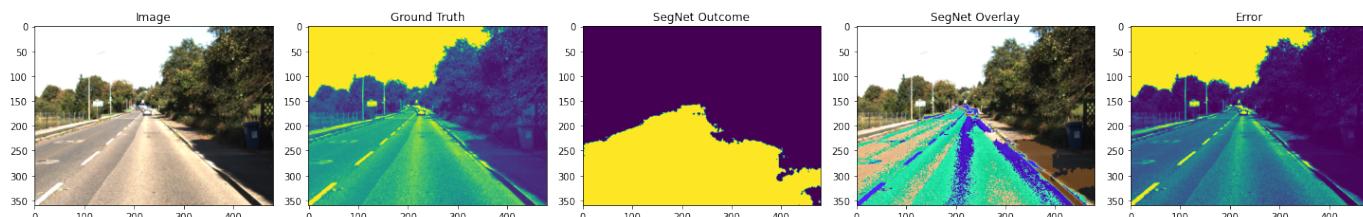
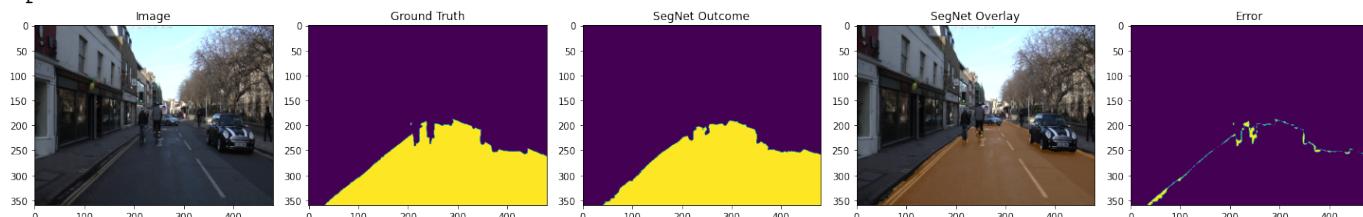
Valid Acc: 98.466%

Time: 85.97s

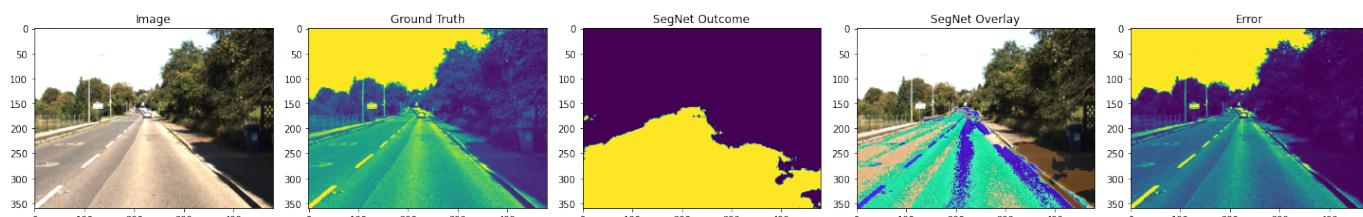
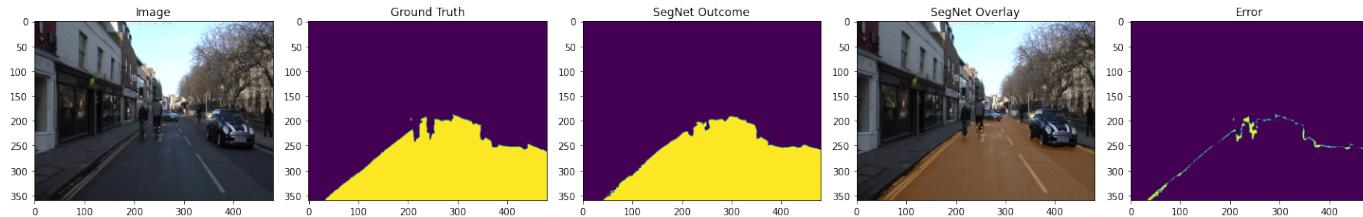




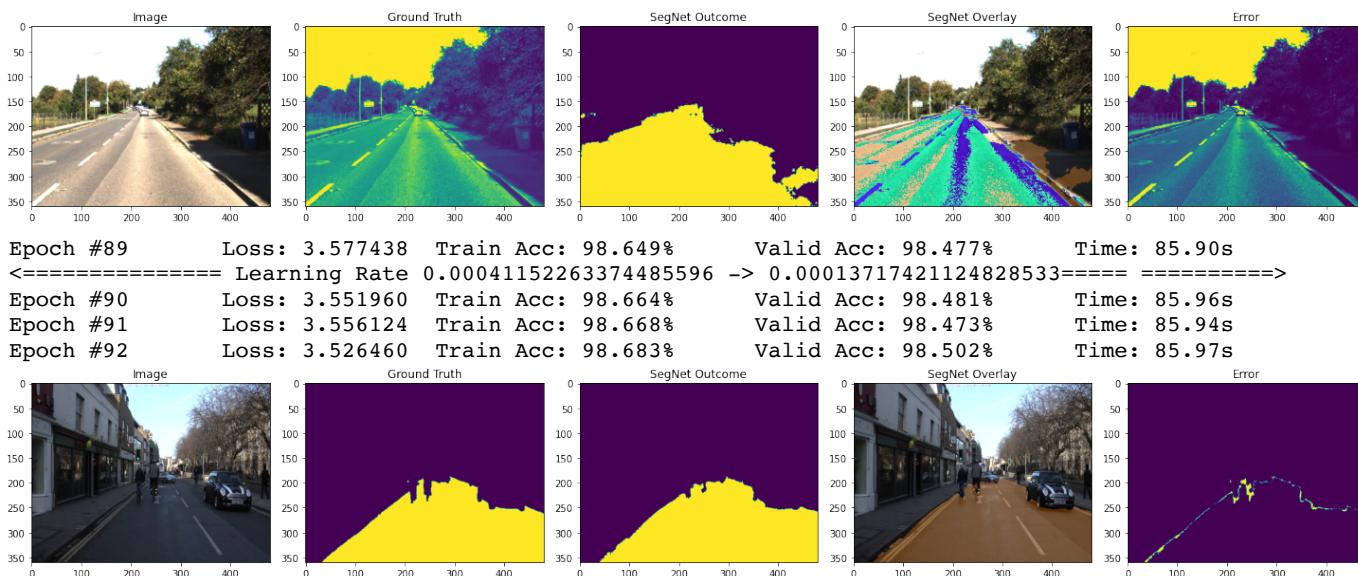
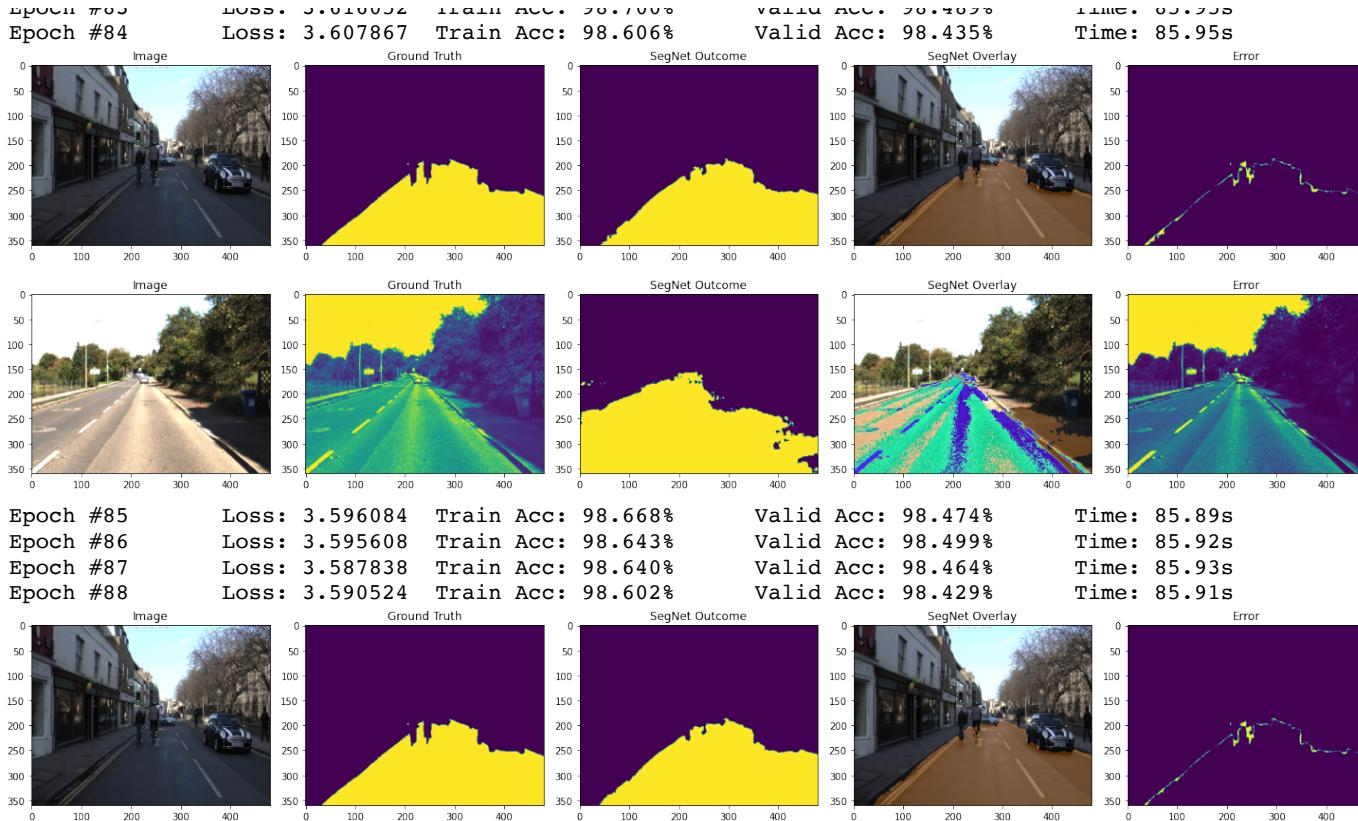
Epoch #73 Loss: 3.864235 Train Acc: 98.503% Valid Acc: 98.438% Time: 85.90s
Epoch #74 Loss: 3.829767 Train Acc: 98.615% Valid Acc: 98.429% Time: 85.94s
<===== Learning Rate 0.0012345679012345679 -> 0.00041152263374485596 =====>
Epoch #75 Loss: 3.708163 Train Acc: 98.633% Valid Acc: 98.481% Time: 85.92s
Epoch #76 Loss: 3.661913 Train Acc: 98.601% Valid Acc: 98.477% Time: 85.92s

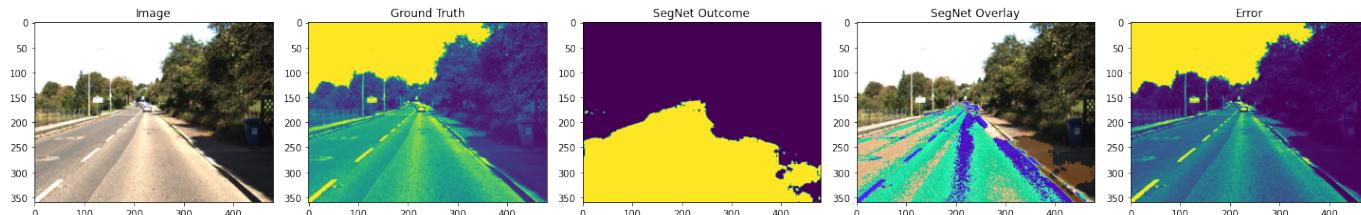


Epoch #77 Loss: 3.662009 Train Acc: 98.628% Valid Acc: 98.486% Time: 85.87s
Epoch #78 Loss: 3.702063 Train Acc: 98.611% Valid Acc: 98.430% Time: 85.92s
Epoch #79 Loss: 3.687805 Train Acc: 98.617% Valid Acc: 98.470% Time: 85.95s
Epoch #80 Loss: 3.657182 Train Acc: 98.630% Valid Acc: 98.487% Time: 85.91s



Epoch #81 Loss: 3.622637 Train Acc: 98.598% Valid Acc: 98.471% Time: 85.94s
Epoch #82 Loss: 3.670632 Train Acc: 98.624% Valid Acc: 98.467% Time: 85.90s
Epoch #83 Loss: 3.616052 Train Acc: 98.700% Valid Acc: 98.499% Time: 85.95s

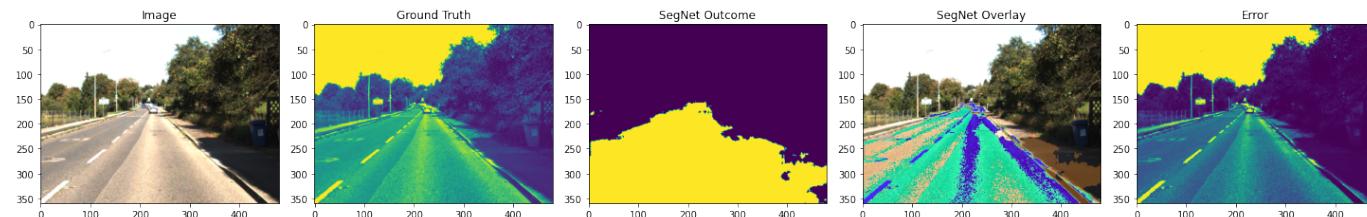
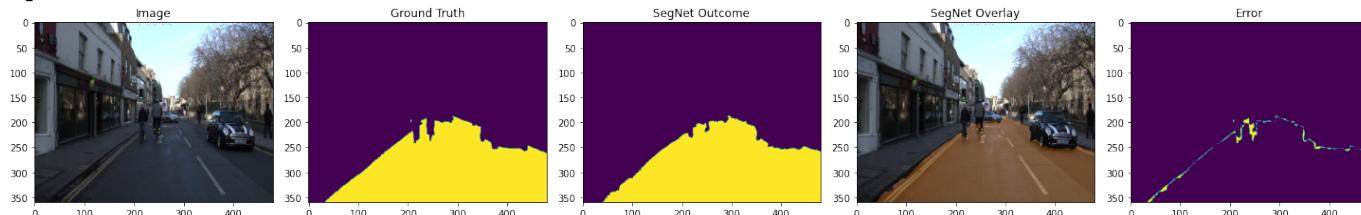




```

Epoch #93      Loss: 3.511168  Train Acc: 98.649%  Valid Acc: 98.459%  Time: 85.95s
Epoch #94      Loss: 3.514572  Train Acc: 98.677%  Valid Acc: 98.482%  Time: 85.94s
Epoch #95      Loss: 3.554174  Train Acc: 98.650%  Valid Acc: 98.459%  Time: 85.93s
Epoch #96      Loss: 3.496509  Train Acc: 98.650%  Valid Acc: 98.469%  Time: 85.92s

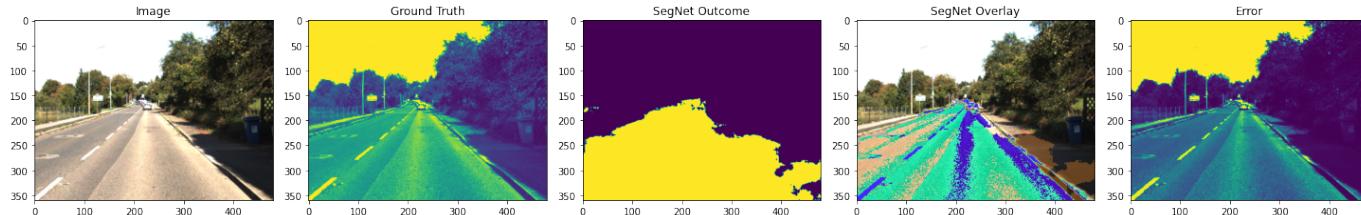
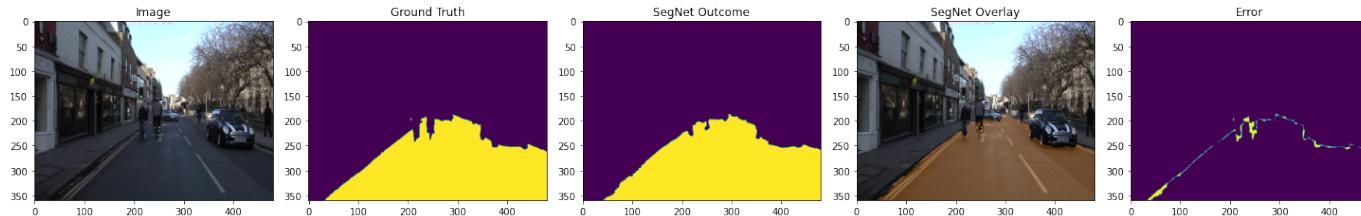
```



```

Epoch #97      Loss: 3.563272  Train Acc: 98.641%  Valid Acc: 98.486%  Time: 85.90s
Epoch #98      Loss: 3.541467  Train Acc: 98.649%  Valid Acc: 98.482%  Time: 85.95s
Epoch #99      Loss: 3.499984  Train Acc: 98.671%  Valid Acc: 98.464%  Time: 85.94s
Epoch #100     Loss: 3.521144  Train Acc: 98.666%  Valid Acc: 98.473%  Time: 85.96s

```



```

Epoch #101     Loss: 3.525872  Train Acc: 98.650%  Valid Acc: 98.460%  Time: 85.93s

```

KeyboardInterrupt

Traceback (most recent call last)

<https://github.com/PyTorchLightning/lightning-bolts>

```
<ipython-input-11-acde988cc0aa> in <module>()
    8
    9 # Launch training
--> 10 train(model, optimizer, loss, X_train, y_train, X_valid, y_valid)

<ipython-input-10-1292fd992415> in train(model, optimizer, loss_fn, X_train, y_train, X_valid, y_valid)
)
    73
    74     # train_acc = avg_pixelwise_accuracy(model, train_dataset)
--> 75     train_acc = avg_pixelwise_accuracy(model, train_dataset) * 100
    76     valid_acc = avg_pixelwise_accuracy(model, valid_dataset) * 100
    77     print("Epoch #{0}\tLoss: {:.6f}\tTrain Acc: {:.3f}%\tValid Acc: {:.3f}%\tTime: {:.2f}s".format(epoch+1, epoch_loss, train_acc, valid_acc, delta))

<ipython-input-5-967846405a9e> in avg_pixelwise_accuracy(model, dataset)
    13
    14     correct, total = 0, 0
--> 15     for i, batch in enumerate(loader):
    16
    17         # mount to GPU if available [Need to fix y]

/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py in __next__(self)
    433         if self._sampler_iter is None:
    434             self._reset()
--> 435         data = self._next_data()
    436         self._num_yielded += 1
    437         if self._dataset_kind == _DatasetKind.Iterable and \


/usr/local/lib/python3.6/dist-packages/torch/utils/data/dataloader.py in _next_data(self)
    473     def _next_data(self):
    474         index = self._next_index() # may raise StopIteration
--> 475         data = self._dataset_fetcher.fetch(index) # may raise StopIteration
    476         if self._pin_memory:
    477             data = _utils.pin_memory.pin_memory(data)

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/fetch.py in fetch(self, possibly_batched_index)
    45
    46         else:
--> 47             return self.collate_fn(data)

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/collate.py in default_collate(batch)
    71     return batch
    72     elif isinstance(elem, container_abcs.Mapping):
--> 73         return {key: default_collate([d[key] for d in batch]) for key in elem}
    74     elif isinstance(elem, tuple) and hasattr(elem, '_fields'): # namedtuple
    75         return elem_type(*(default_collate(samples) for samples in zip(*batch)))

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/collate.py in <dictcomp>(.0)
    71     return batch
    72     elif isinstance(elem, container_abcs.Mapping):
--> 73         return {key: default_collate([d[key] for d in batch]) for key in elem}
    74     elif isinstance(elem, tuple) and hasattr(elem, '_fields'): # namedtuple
    75         return elem_type(*(default_collate(samples) for samples in zip(*batch)))

/usr/local/lib/python3.6/dist-packages/torch/utils/data/_utils/collate.py in default_collate(batch)
    53         storage = elem.storage().__new_shared(numel)
    54         out = elem.new(storage)
```

```
---> 55         return torch.stack(batch, 0, out=out)
56     elif elem_type.__module__ == 'numpy' and elem_type.__name__ != 'str_' \
57         and elem_type.__name__ != 'string_':
```

KeyboardInterrupt:

Train from VGG weights

In []:

```
TRANSFER_LEARNING = True
model_vgginit = SegNet(3, 2, transfer_learning=TRANSFER_LEARNING).to(DEVICE)
optimizer = torch.optim.Adam(model_vgginit.parameters(), lr=LEARNING_RATE)
loss = nn.CrossEntropyLoss(weight=LABEL_WEIGHTS).to(DEVICE)

# Launch training
torch.cuda.empty_cache()
train(model_vgginit, optimizer, loss, X_train, y_train, X_valid, y_valid)
```

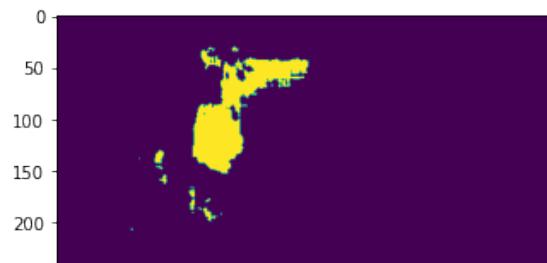
Load model

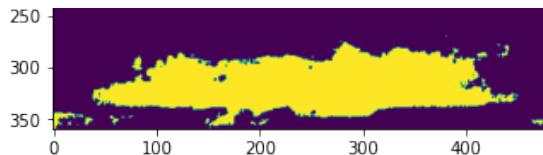
In [9]:

```
# model_name = r'Epoch101_loss3.5258_trainacc98.650_valacc98.460.pth' # augmented
# model_name = r'../Epoch49_loss2.1120_trainacc97.727_valacc97.859.pth' # scratch
model_name = r'../Epoch49_loss2.1120_trainacc97.727_valacc97.859.pth' # scratch
model_load = load_checkpoints(model_name)

show_pred_mask(model_load, X_train[0, :])
model = model_load
```

Model Loaded: Epoch#49 Loss:2.1120 TrainAcc:97.7268 ValidAcc:97.8592





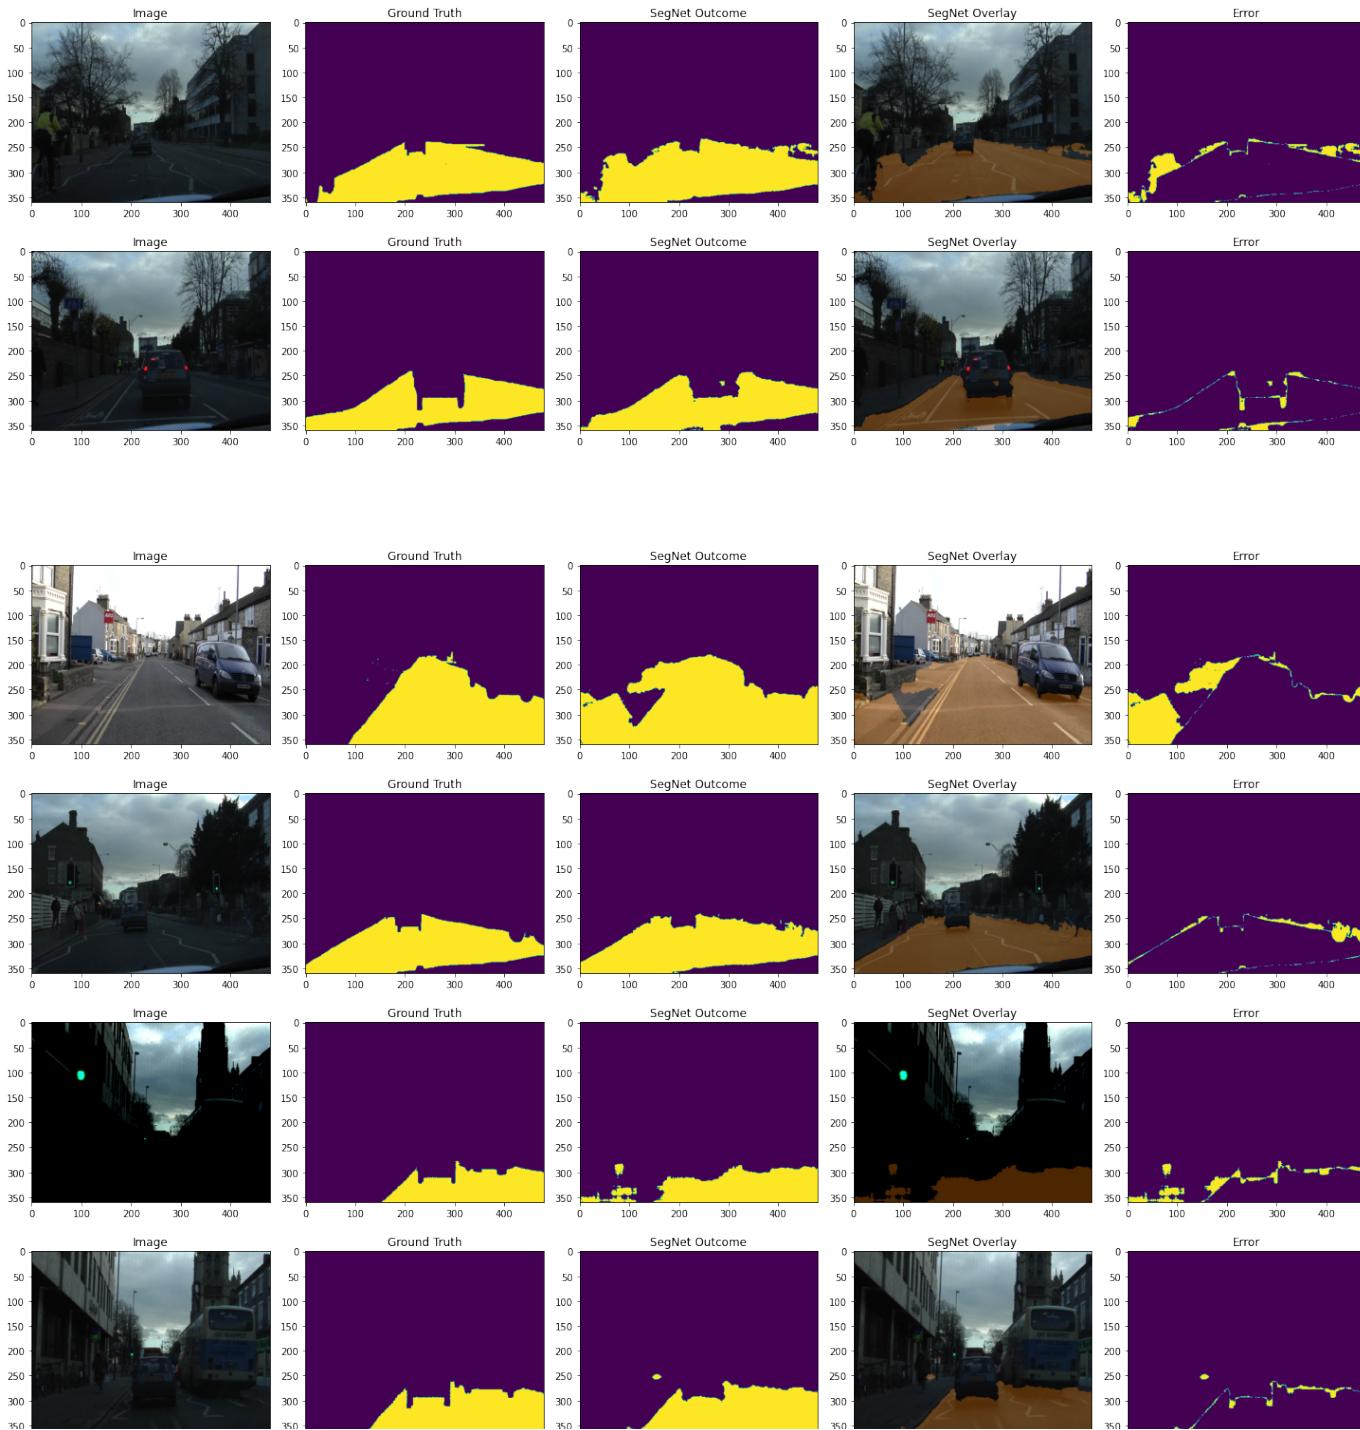
Report Evaluation results on test dataset

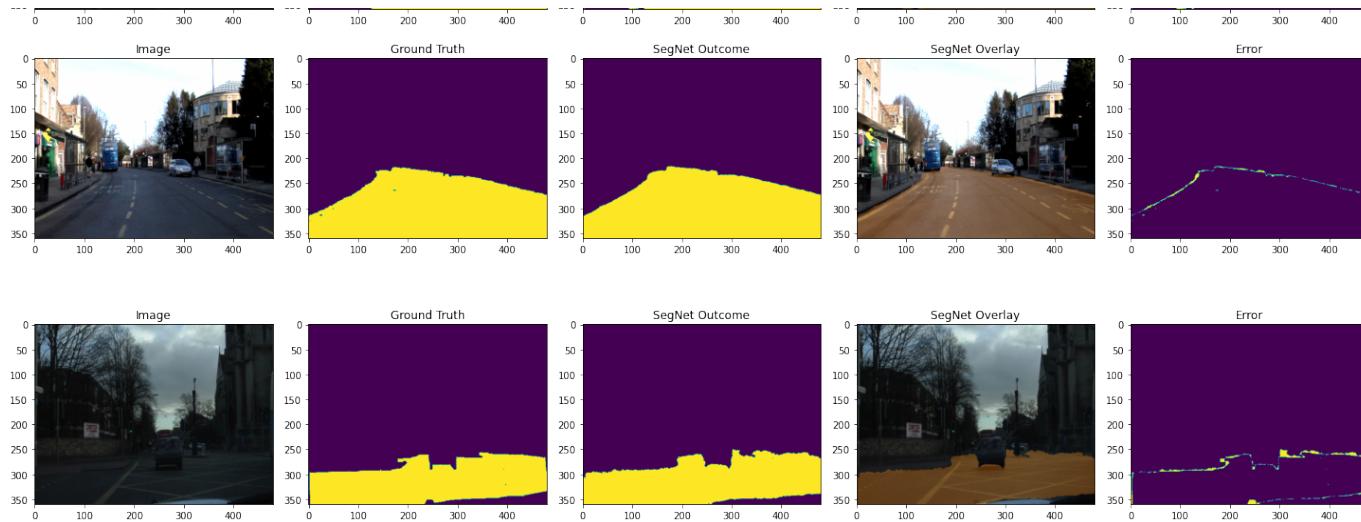
```
In [10]:  
torch.cuda.empty_cache()  
  
model_selected = model  
test_dataset = data.CamVidDataset(X_test, y_test)  
test_accuracy = avg_pixelwise_accuracy(model_selected, test_dataset)  
print("Selected model's pixelwise accuracy on test dataset : {:.5f}%".format(test_accuracy * 100))  
  
Selected model's pixelwise accuracy on test dataset : 90.24887%
```

Visualization Output

```
In [15]:  
# visualize a single input image  
# show_img(X_test[0, :])  
# show_mask(y_test[0, 1, :])  
# show_pred_mask(model_load, X_test[0, :])  
  
# cmap = 'prism'  
# cmap = 'brg'  
# cmap = 'Paired'  
# cmap = 'jet'  
  
model_selected = model  
show_all(model_selected, X_test[2, :], y_test[2, :], cmap='jet')  
show_all(model_selected, X_test[11, :], y_test[11, :], cmap='jet')  
show_all(model_selected, X_test[-2, :], y_test[-2, :], cmap='jet')  
show_all(model_selected, X_test[43, :], y_test[43, :], cmap='jet')  
  
show_all(model_selected, X_train[2, :], y_train[2, :], cmap='jet')  
show_all(model_selected, X_train[11, :], y_train[11, :], cmap='jet')  
show_all(model_selected, X_train[-2, :], y_train[-2, :], cmap='jet')
```

```
snow_all(model_selected, x_train[45, :, :], y_train[45, :, :], cmap= jet )
```





Demostration

On Demo.py

```
In [16]: model.eval()

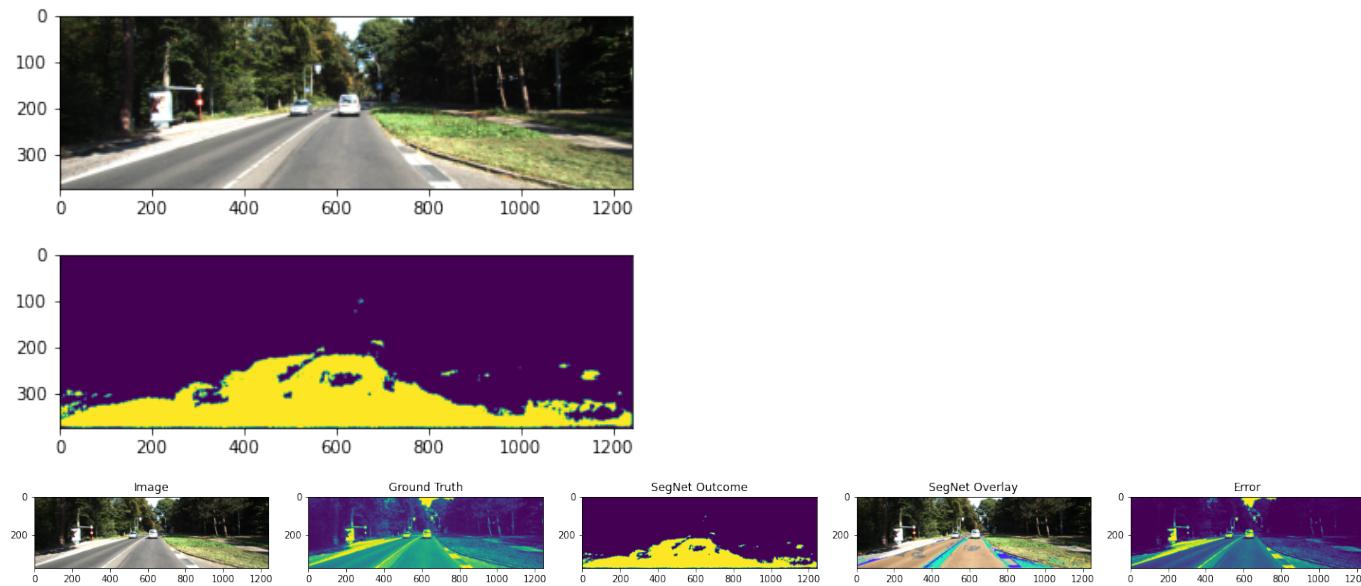
kitti_img = cv2.imread("/content/gdrive/MyDrive/Segmentation/Kitti/um_00007.png")

kitti_tensor = torch.from_numpy(np.moveaxis(kitti_img, -1, 0))

show_img(kitti_tensor)

show_pred_mask(model, kitti_tensor.float())

show_all(model, kitti_tensor.float(), kitti_tensor.float()[0, :])
```



Overlay Segmentation results

```
In [18]: def write_to_dir(model, load_dir, save_dir):

    # change to eval mode
    model.eval()

    # create output dir
    os.makedirs(save_dir, exist_ok=True)

    # loop over the given directory
    for item in sorted(os.listdir(load_dir)):
```

```
for item in os.listdir(os.path.join(load_dir, item)):
    if '.png' in item:
        img = cv2.imread(os.path.join(load_dir, item))
        img = cv2.resize(img, (480, 320))
        img = np.moveaxis(img, -1, 0)
        img = torch.from_numpy(img).type(torch.float32)
        out = model(img.unsqueeze(0).to(DEVICE))
        predicted = torch.argmax(out, dim=1)
        predicted = predicted.cpu().data.numpy()[0, :]

        img = torch.movedim(img, 0, -1).type(torch.uint8).numpy().astype(np.float)
        img[:, :, 1] += (predicted * 255 * 0.5).astype(np.uint8)
        img = np.clip(img, 0, 255).astype(np.uint8)
        # img[:, :, 0] += (predicted * 255 * 0.7).astype(np.uint8)

        cv2.imwrite(os.path.join(save_dir, item), img)
```

```
In [19]:  
# model_name = r'Epoch20_loss1.5263_trainacc97.227_valacc97.571.pth'  
model_load = model  
  
load_dir = r'/content/gdrive/MyDrive/KiTTi_dataset/testing_crop'  
save_dir = r'/content/gdrive/MyDrive/KiTTi_dataset/testing_crop_segnet_new'  
write_to_dir(model_load, load_dir, save_dir)
```

```
In [21]:  
load_dir = r'/content/gdrive/MyDrive/KiTTi_dataset/training_crop'  
save_dir = r'/content/gdrive/MyDrive/KiTTi_dataset/training_crop_segnet_new'  
write_to_dir(model_load, load_dir, save_dir)
```

```
In [22]:  
load_dir = camvid_path+"/train"  
save_dir = camvid_path+"/train_segnet_new"  
write_to_dir(model_load, load_dir, save_dir)  
  
load_dir = camvid_path+"/test"  
save_dir = camvid_path+"/test_segnet_new"  
write_to_dir(model_load, load_dir, save_dir)  
  
load_dir = camvid_path+"/val"  
save_dir = camvid_path+"/val_segnet_new"  
write_to_dir(model_load, load_dir, save_dir)
```